

Below is a rewritten version of the PAQJP 9.2 project portfolio. The content has been reorganised for clarity, condensed where appropriate, and polished for academic presentation – while preserving all essential technical details, formulas, and experimental results.

PAQJP 9.2: Universal Lossless Compression via 65 536 Invertible Transformation Paths

Final Project Portfolio – Data Science

Written by Jurijus Pacalovas

Abstract

PAQJP 9.2 reformulates lossless compression as an exhaustive optimisation problem over a space of 65 536 mathematically invertible transformation paths. A library of 256 bijective byte-level operations – ranging from run-length coding and bit-rotation to number-theoretic Fermat-Little-Theorem mappings and canonical Huffman encoding – provides a reversible feature-engineering layer. By considering every ordered pair of two such transforms (excluding the double identity) and adding an explicit raw-data path, exactly 65 536 distinct preprocessing pipelines are obtained.

The compressor performs a brute-force search over all 65 536 paths: it applies each candidate transform (or pair) to the input, compresses the result with a backend (Zstandard, PAQ, or identity), and verifies perfect bit-exact reconstruction before acceptance. The chosen path is encoded in a compact variable-length header (1–3 bytes) using a marker-based scheme, and the payload is the pure compressed stream with no internal markers. Decompression is deterministic and requires no trial-and-error.

A built-in self-test verifies the invertibility of every single transform and every composite pair, guaranteeing that the system will never emit a corrupted file. This portfolio presents

the mathematical structure, the optimisation framework, and experimental evidence for the PAQJP 9.2 system.

1. Introduction

Standard lossless compression tools (Zstandard, PAQ) operate directly on raw byte sequences, exploiting statistical dependencies. PAQJP 9.2 introduces a reversible transformation stage before compression. The central observation is that many bijective re-encodings of the input can expose additional redundancy, allowing the backend compressor to achieve smaller outputs.

Compression is thus recast as a supervised model-selection problem:

- Feature engineering – apply an invertible byte-level transform (or an ordered pair of transforms) to the data.
- Model training – compress the transformed stream with an off-the-shelf codec.
- Validation – decompress and compare bit-for-bit with the original.
- Hyper-parameter tuning – exhaustive search over 256 single transforms, 65 535 non-trivial ordered pairs, and the raw identity (total 65 536 candidate pipelines).

The entire system is provably lossless because every elementary transform is a bijection, and the compressor rejects any output that does not verify.

2. Mathematical Foundations

2.1 Notation

Let $\mathbb{B} = \{0,1,\dots,255\}$ be the set of byte values.

Let $B = (B[0], B[1], \dots, B[n-1]) \in \mathbb{B}^n$ denote the input byte sequence.

We use the following operators:

- $\oplus$ – bitwise XOR,
- $\ll, \gg$ – left/right bit shift,
- $\text{rotl}(x, r)$ – cyclic left rotation of an 8-bit value by r bits,
- $\bmod$ – modulo operation.

2.2 Transform Definition

A transform is a bijective function $T : \mathbb{B}^n \rightarrow \mathbb{B}^m$ together with an explicit inverse T^{-1} such that for every input B ,

$$T^{-1}(\text{bigl}(T(B)\bigr)) = B.$$

The output length m may differ from n only for a few transforms (run-length encoding, prefix-code compression) that store auxiliary parameters; invertibility is always preserved.

3. The 256 Single Transforms – Formula Families

The 256 elementary transforms are organised into four broad families. Representative formulas are presented below; the complete set follows analogous bijective constructions.

3.1 Original Set (Transforms 1–24)

These include XOR-based, arithmetic, substitution, rotation, run-length, and constant-mask transforms.

Transform 01 – Prime-XOR

For each prime $p < 256$, compute $x_p = \max(1, \lceil p \cdot 4096 / 28672 \rceil)$. Repeated R times, for every index i with $i \bmod 3 = 0$:

$$B[i] \leftarrow B[i] \oplus x_p.$$

Inverse: identical (XOR is self-inverse).

Transform 04 – Positional Subtraction

Forward (repeated R times):

$$C[i] = \text{bigl}(B[i] - (i \bmod 256) \text{bigr} \bmod 256.$$

Inverse:

$$B[i] = \text{bigl}(C[i] + (i \bmod 256) \text{bigr} \bmod 256.$$

Transform 06 – Substitution Cipher

A Fisher-Yates permutation $\pi : \mathbb{B} \rightarrow \mathbb{B}$ is generated with seed 42. Then

$$C[i] = \pi(B[i]), \text{quad } B[i] = \pi^{-1}(C[i]).$$

Transform 05 – Bit Rotation

Fixed 3-bit left rotation:

$$C[i] = \text{rotl}(B[i], 3), \quad B[i] = \text{rotr}(C[i], 3).$$

Transform 00 – Multi-Pass RLE

An adaptive run-length encoder: each pass finds an additive shift $s \in [0, 255]$ that maximises the sum of squared run lengths. The shifted data is encoded with a variable-length prefix code. Header: number of passes + shifts. Inverse: decode runs, then apply reverse shifts in opposite order.

Transform 14 – Checksum Append

$$C = B \bigoplus_{i=0}^{n-1} B[i].$$

Inverse: drop the last byte.

Transform 23 – Iterative Constant Diapason

Each 4-bit nibble is mapped to a shorter bit sequence using a fixed prefix-free code. The mapping is applied repeatedly until the total bit length stops decreasing. Original bit length and pass count are stored. Inverse: repeatedly decode the prefix codes and reassemble bits.

Transform 24 – 43-Byte Block RLE

Data is processed in 43-byte blocks. If all bytes in a block are equal, a flag bit (1) is followed by the byte value and a 6-bit (length-1). Otherwise, flag bit (0) followed by 6-bit length and the raw bytes.

3.2 Fermat-Little-Theorem Transforms (Transforms 25–30)

These transforms exploit the bijection on $\mathbb{B}$ given by

$$f_e(x) = \text{bigl}((x+1)^e \bmod 257 \bigr) - 1 \bmod 256 ,$$

where e is odd and $\gcd(e, 256)=1$. The inverse exponent d satisfies $e \cdot d \equiv 1 \pmod{256}$; then

$$f_e^{-1}(y) = \text{bigl}((y+1)^d \bmod 257 \bigr) - 1 \bmod 256 .$$

- Transform 25 – fixed $e = 3$ ($d = 171$).
- Transform 26 – two-byte parameter n derived from length; $e = (n^{16} \bmod 216) \bmod 256 \mid 1$.
- Transform 27 – 1024-byte blocks; each block uses e^{200} with block-specific n .
- Transform 28 – same as 27, but the transformed block is further compressed by the backend.
- Transform 29 – 32-byte blocks; exponent always 1 (identity), only backend compression applied.
- Transform 30 – 33-byte blocks; n is either computed from the first byte or, if out of range, from SHA-256 of the block; identity Fermat, only backend compression.

All Fermat transforms are provably bijective.

3.3 Dynamic XOR Transforms (Transforms 31–40 and 48–255)

For transform index t , let

$$\text{seed} = \text{SeedTable}[t \bmod 126]; [\text{length} \bmod 40];$$

Then

$C[i] = B[i] \oplus \text{seed}$.

Inverse is identical. The seed table is a fixed 126×40 matrix of pseudo-random bytes, making the XOR masks unpredictable without the table.

3.4 Special Named Algorithms (Transforms 41–47)

- Transform 41 – Algorithm 2703 – XOR the first 8 bytes with the pair $\{0x27, 0x03\}$ cyclically.
- Transform 42 – Extended 2703 – XOR every byte with $\{0x27, 0x03\}$ cyclically.
- Transform 43 – Base 16 777 216 – XOR 3-byte blocks with the mask $\{0x10, 0x00, 0x00\}$.
- Transform 44 – Base64 – Encode data in Base64; decode reverses the operation.
- Transform 45 – Canonical Huffman – Build a Huffman tree from byte frequencies, encode with canonical codes. Header stores original length and 256 code lengths. Inverse: rebuild the canonical code table and decode.
- Transform 46 – Power-of-2/3/6 Minus-10 XOR – A 100-byte mask, constructed from the sequence $[1, 2, 4, 8, 16, 32, 64, 128, 3, 6]$ with 10 subtracted from each element, repeated 10 times, is XORed cyclically.
- Transform 47 – PAQ State Table XOR – A modified version of the PAQ state table (each entry minus 400, masked to a byte) is used: for each byte b , its row $b \bmod 256$ is looked up, and the first column value is XORed with b . Self-inverse.

3.5 Identity (Transform 256)

$T_{256}(B) = B$.

4. Transform Pairs and the 65 536-Path Search Space

By chaining two transforms we can achieve more powerful rearrangements. For ordered pair (a,b) with $1 \leq a,b \leq 256$, the composite is

$$T_{\{(a,b)\}}(B) = T_b \circ T_a(B).$$

The inverse is

$$T_{\{(a,b)\}}^{-1}(C) = T_a^{-1} \circ T_b^{-1}(C).$$

Because the identity $T_{\{256\}}$ exists, every single transform is a special case: for any t , the pair $(t,256)$ applies exactly T_t . There are $256 \times 256 = 65,536$ ordered pairs. The pair $(256,256)$ is the double identity, equivalent to no transformation. We replace it with an explicit raw path, so the candidate set becomes

$$\Theta = \{\text{raw}\} \cup \{T_{\{(a,b)\}}; 1 \leq a,b \leq 256, (a,b) \neq (256,256)\}.$$

Hence $|\Theta| = 1 + (65,536 - 1) = 65,536$ distinct transformation paths.

Path Index Encoding

For a pair (a,b) the linear index is

$$p = (a-1) \times 256 + (b-1).$$

This maps $(1,1)$ to 0 and $(256,256)$ to 65,535. We exclude $p = 65,535$ and collect the remaining 65,535 pairs in increasing order of p . We then assign:

- Path index 0 → raw (no transform).
- Path index i ($1 \leq i \leq 65535$) → the pair whose i is the i -th smallest among the non-excluded indices.

Thus every candidate pipeline is identified by a single integer $i \in \{0, \dots, 65535\}$.

5. Header Encoding

The chosen path index is stored in a compact, marker-free header that precedes the compressed payload. The scheme uses escape bytes to distinguish four cases:

Path type Header bytes

Single transform 1–252 [t-1]

Single transform 253–256 [254, t-253]

Raw data (index 0) [252]

Transform pair (indices 1–65535) [253, $(p \gg 8) \& 0xFF$, $p \& 0xFF$] where p is the pair's linear index (0...65 534)

The maximum header length is 3 bytes. No special markers are inserted into the payload; the compressed stream produced by the backend is directly appended.

Decompression reads the first byte:

- if < 252 → apply single transform $f+1$;
- if $= 252$ → raw;
- if $= 253$ → read two more bytes, recover p , map to (a,b) , and apply the inverse pair;

- if $=254 \rightarrow$ read one more byte x , apply transform $253+x$.

The reverse transforms are applied after the backend decompressor restores the transformed data.

6. Compression Backend and Marker-Free Decompression

The backend compressor $\mathcal{C}$ is chosen as the smallest result among:

- Zstandard (level 22) if available,
- PAQ if available,
- raw byte storage (identity).

No type marker is prepended to the payload. During decompression, the raw payload is fed to the backends in a fixed order: first try Zstandard, then PAQ, and finally assume uncompressed data. The first backend that succeeds and produces a valid byte stream (as verified during compression) yields the transformed data. This is safe because the compression phase already validated every accepted candidate: only byte sequences that are unambiguous and correctly decode are kept.

7. Optimisation Objective and Lossless Verification

For a given input B , the compressor evaluates every path index $i \in \{0, \dots, 65, 535\}$:

1. Compute $X_i = T_i(B)$ (using the transform or pair associated with i ; T_0 is the raw identity).

2. Compute $Z_i = \mathcal{C}(X_i)$.
3. Form the full compressed stream $S_i = \text{Header}(i) \backslash; \backslash; Z_i$.
4. Verify: decompress S_i using the marker-free decoder and compare the result with B. If they are not bit-identical, reject candidate i .
5. Select

$$i^* = \arg\min_{\{i \in \{0, \dots, 65535\} \mid \text{verified}\}} \text{len}(S_i).$$

The final output is S_{i^*} . If no candidate passes verification, the compressor refuses to produce a file – a hard guarantee against data corruption.

Losslessness Theorem

Since every T_i is a composition of bijective elementary transforms, $\mathcal{C}$ is chosen losslessly, and the verification step is mandatory, the pipeline is provably lossless. All mathematical components are invertible, and the verification catches any implementation error.

8. Self-Test – Exhaustive Invertibility Check

To confirm implementational correctness, an exhaustive self-test is performed:

- For every path index $i \in \{0, \dots, 65535\}$, the system applies the forward transformation to a fixed test byte (e.g., 0xAA), then applies the inverse, and checks that the original byte is recovered.
- For a 1000-byte random block, a full compression–decompression cycle is run, and bit-exact identity is verified.

The test passes on all 65 536 paths, demonstrating the mathematical invertibility of every component and the correctness of the composite mappings.

9. Experimental Results

The framework was evaluated on a variety of real-world files. The compression ratio $\text{CR} = \frac{\text{compressed size}}{\text{original size}}$ is reported.

File type	Backend	Best path	index	CR (backend alone)	CR (PAQJP 9.2)
-----------	---------	-----------	-------	--------------------	----------------

JPEG	(24 KB)	PAQ raw	(0)	0.964	0.964
------	---------	---------	-----	-------	-------

PNG screenshot	(52 KB)	PAQ RLE-pair	(non-zero)	1.000 (raw fallback)	0.745
----------------	---------	--------------	------------	----------------------	-------

English text	(100 KB)	PAQ pair	(2,17)	0.234	0.187
--------------	----------	----------	--------	-------	-------

PDF	(200 KB)	Zstd raw	(0)	0.987	0.987
-----	----------	----------	-----	-------	-------

The most striking improvement appears on text, where a carefully chosen transform pair reduced the PAQ-only size by an additional 20%. On a PNG screenshot, a pair containing an RLE transform turned a previously incompressible file into one that compressed to 74.5% of its original size. In all tests, no file ever caused a false acceptance; the compressor refused output when ambiguity could arise, thereby preserving perfect losslessness.

10. Conclusion

PAQJP 9.2 demonstrates that a data-driven approach – building a large space of mathematically invertible byte-level transforms and exhaustively searching it – can significantly improve compression ratios over standard algorithms. Every component is defined by an explicit formula, every transform is proven bijective, and the marker-free design with built-in self-verification guarantees zero-error decompression.

The system integrates feature engineering, model selection, and rigorous validation into a single pipeline. The 65 536-path search space, enriched by number-theoretic, statistical, and structural transforms, offers a flexible, general-purpose preprocessing layer that adapts to any input, making PAQJP 9.2 a substantial contribution to lossless compression and a notable data-science project.

All mathematical formulas are contained in this presentation. Full source code (Python) is available upon request.