

PJP-6: A Universal Lossless Compression Framework with 2,961 Transform Paths and Hybrid Dictionary Support

Abstract

PJP-6 is an advanced lossless compression system that reformulates compression as a model-selection problem over a large space of mathematically invertible byte-level preprocessing pipelines. Building upon the principles of PAQJP, PJP-6 extends the original design with hybrid dictionary compression, quantum-inspired permutations, and specialised transforms for structured data such as DOCX documents. The system comprises 256 elementary bijective transforms, from which a curated subset of 52 safe transforms is used to generate 2,704 ordered pairs. Together with the single transforms and the raw identity, this yields exactly 2,961 distinct transformation sequences – plus additional quantum and dictionary-aware modes. An exhaustive search evaluates every candidate pipeline, compresses the transformed data with either Zstandard, PAQ, or raw storage, and verifies bit-exact reconstruction before acceptance. A compact variable-length header (1–3 bytes) stores the chosen path. The decompressor is deterministic and requires no trial-and-error. Built-in self-tests confirm the invertibility of all transforms on every byte value, guaranteeing that the system never emits corrupted data. This document describes the mathematical foundations, the implementation architecture, and the experimental validation of PJP-6.

1. Introduction

Standard lossless compression tools such as Zstandard and PAQ operate directly on raw byte sequences, exploiting statistical redundancies. PJP-6 introduces a reversible transformation stage prior to compression. The key insight is that many bijective re-encodings of the input can expose additional redundancy, allowing the backend compressor to produce smaller outputs.

Compression is thus cast as a supervised learning pipeline:

- Feature engineering – apply an invertible byte-level transform (or an ordered pair) to the data.
- Model training – compress the transformed stream with an off-the-shelf codec.
- Validation – decompress and compare bit-for-bit with the original.
- Hyper-parameter tuning – exhaustive search over 256 single transforms, 2,704 non-trivial ordered pairs, the raw identity, and optionally dictionary-based and quantum paths.

Every elementary transform is a mathematical bijection, and the compressor rejects any output that does not verify. The entire system is provably lossless. PJP-6 goes beyond the original framework by incorporating hybrid dictionary compression (static word, line, and dynamic tokenisation), Qiskit-inspired quantum permutations, and specialised transforms for DOCX paragraphs and tables. This expanded feature set enables superior compression ratios on text, structured documents, and certain binary formats while preserving the core guarantee of zero-error decompression.

2. Mathematical Foundations

2.1 Notation

Let $\mathbb{B} = \{0, 1, \dots, 255\}$ be the set of all byte values. A byte sequence of length n is denoted as $\mathbf{B} = (B[0], B[1], \dots, B[n-1]) \in \mathbb{B}^n$. Operators include bitwise XOR ($\oplus$), left/right shifts ($\ll, \gg$), cyclic rotations (rotl , rotr), and modulo arithmetic ($\bmod$). A transform is a bijective function $T : \mathbb{B}^n \rightarrow \mathbb{B}^n$ with an explicit inverse T^{-1} such that for every input $\mathbf{B}$,

$$T^{-1}(T(\mathbf{B})) = \mathbf{B}.$$

Composition of transforms yields new bijections; the inverse of a composition is the reverse composition of the inverses.

2.2 The 256 Single Transforms

The elementary transforms are organised into several families, each designed to be bijective on bytes. Below are the key ones with their formal definitions.

Constant XOR with a prime-based value (Transform 2,9,14). For each prime $p < 256$ and a parameter r ,

$$x_p = \max\{1, \lceil p \cdot 4096 / 28672 \rceil\}.$$

Then, for every third byte, XOR with x_p repeated r times. The transform is self-inverse because XOR is its own inverse.

Index-based arithmetic (Transform 5). Each byte is subtracted by its positional index modulo 256:

$$T_5(\mathbf{B})[i] = (B[i] - i) \bmod 256.$$

The inverse adds back the index:

$$T_5^{-1}(\mathbf{C})[i] = (C[i] + i) \bmod 256.$$

Fixed rotation (Transform 6). A 3-bit left rotation applied to every byte:

$$T_6(B[i]) = \text{rotl}(B[i], 3) = \text{bigl}((B[i] \ll 3) \vee (B[i] \gg 5)\text{bigr} \wedge 255.$$

The inverse uses a right rotation by 3.

Substitution cipher (Transform 7). A Fisher-Yates permutation π on $\{0, \dots, 255\}$ with a fixed seed:

$$T_7(B[i]) = \pi(B[i]), \quad T_7^{-1}(C[i]) = \pi^{-1}(C[i]).$$

Constant shift (Transform 22). Adding 255 (i.e., subtracting 1) modulo 256:

$$T_{22}(B[i]) = (B[i] + 255) \bmod 256, \quad T_{22}^{-1}(C[i]) = (C[i] - 255) \bmod 256.$$

PI- and constant-based mask transforms (Transforms 17–21). Using decimal expansions of π , $\pi^2/6$, $1/e$, $5e$ to create cyclic XOR masks M of length L :

$$T_k(B[i]) = B[i] \oplus M[i \bmod L], \quad \text{self-inverse}.$$

Dynamic XOR transforms (Transforms 23–255). For transform index t , a seed s is fetched from a precomputed table using t and data length; then

$$T_t(B[i]) = B[i] \oplus s, \quad \text{self-inverse}.$$

Identity (Transform 256). $T_{256}(\mathbf{B}) = \mathbf{B}$.

3. Transform Pairs and the Search Space

By chaining two transforms we achieve more powerful rearrangements. A curated subset of the first 52 safe (bijective) transforms is used to form ordered pairs. For an ordered pair (a, b) , the composite is

$$T_{\{(a,b)\}}(\mathbf{B}) = T_b \circ (T_a(\mathbf{B}))$$

Its inverse is

$$T_{\{(a,b)\}}^{-1}(\mathbf{C}) = T_a^{-1} \circ (T_b^{-1}(\mathbf{C}))$$

Because the identity $T_{\{256\}}$ exists, every single transform is a special case. The candidate set becomes

$$\Theta = \{\text{raw}\} \cup \{T_t \mid 1 \leq t \leq 256\} \cup \{T_{\{(a,b)\}} \mid 1 \leq a, b \leq 52\}.$$

Thus

$$|\Theta| = 1 + 256 + 52^2 = 1 + 256 + 2704 = 2961.$$

Pairs are indexed linearly: $\text{idx} = (a-1) \times 52 + (b-1)$, a number in $[0, 2703]$. This index is stored in the header when a pair is selected.

4. Header Encoding

The chosen path is stored in a compact, marker-free header that precedes the compressed payload:

- Single transform 1–252: 1 byte = $t-1$.
- Single transform 253–256: 2 bytes = 254, $t-253$.
- Raw data: 1 byte = 252.
- Transform pair (indices 0–2703): 3 bytes = 253, $(\text{idx} \gg 8) \text{ \& } 255$, $\text{idx} \text{ \& } 255$.

The maximum header length is 3 bytes. No special markers are inserted into the payload; the compressed stream from the backend is directly appended. Decompression reads the first byte and dispatches accordingly.

5. Compression Backends – Dual Mode with Auto-Correction

PJP-6 uses a dual-mode backend strategy that guarantees losslessness without sacrificing compression ratio.

Marker-Free Mode (Default). The backend compressor outputs the shortest result among Zstandard (level 22), PAQ, and raw storage (prefixed with a single byte N for identification). No type marker is prepended. During decompression, the payload is fed to Zstandard, then PAQ, and finally raw-if- N ; the first successful decoder yields the transformed data. This mode eliminates per-file overhead but carries a negligible risk of ambiguity – a valid Zstandard stream might by chance begin with the byte N .

Safe Mode (Automatic Fallback). When marker-free mode produces an ambiguous stream (detected by a failed round-trip verification), the compressor automatically re-compresses in safe mode. In safe mode, each backend prepends a one-byte marker: Z for Zstandard, P for PAQ, N for raw. The decompressor reads the first byte and dispatches to the correct engine, guaranteeing unambiguous decoding under all circumstances.

6. Hybrid Dictionary Compression

PJP-6 integrates three dictionary-based compression modes for text and structured data:

- Static Word Dictionary. A large pre-downloaded dictionary maps frequent words to integer indices. The tokeniser splits text into words and separators; known words are replaced by their indices, while unknown words and separators are stored as raw bytes with length prefixes. The dictionary itself is stored in the compressed stream.
- Line Dictionary. The same merged dictionary is used to find the longest matching phrase; tokenisation proceeds greedily, replacing long common phrases with indices. This is particularly effective for repetitive text such as log files or source code.
- Dynamic Dictionary. This mode builds a frequency-based dictionary on-the-fly from the input text. The text is split into chunks (paragraphs, lines, sentences, words), and the most frequent chunks are assigned short codes. The dictionary is serialised into the output, enabling effective compression of any text without a pre-existing corpus.

Each dictionary mode prepends a unique magic number to the compressed stream, so the decompressor can identify and reverse the appropriate tokenisation.

7. Quantum-Inspired Transforms

Optionally, PJP-6 can employ quantum-inspired transforms generated using Qiskit. A quantum circuit with a given number of qubits (8 for fast, 12 for ultra) is constructed with random rotations and CNOT gates. The circuit's QASM string is hashed to seed a random permutation of the byte values (for 8-bit) or a permutation of 2704-byte blocks (for ultra). These permutations are used as substitution or block-level permutations, effectively introducing non-linear, hard-to-predict transformations that may expose hidden patterns. The system includes 9 fast transforms (8-bit permutations) and 17 ultra transforms (2704-block permutations), all of which are bijective and verified through self-test. Formally, for a permutation σ on the set of all 2^8 byte values, the forward transform is:

$$T_{\{\sigma\}}(\mathbf{B})[i] = \sigma(B[i]).$$

The inverse uses σ^{-1} .

8. Additional Transforms for Structured Data

PJP-6 introduces specialised transforms for specific data types:

- 6-bit Text Compression (Transform 27). For text that uses only a 64-character alphabet (A-Z, a-z, 0-9, space, newline), each character is packed into 6 bits, achieving a 25% reduction. The transform is lossless only for input that conforms to the alphabet; otherwise it falls back to raw. Let Σ be the 64-character set, and $f: \Sigma \rightarrow \{0, \dots, 63\}$ be the bijection. The packed bit-length is $\lceil |\text{text}\{text\}| \rceil \times 6 \lceil \rceil$ bytes.

- Per-3-Byte Subtraction (Transforms 28–30). These transforms group the data into 3-byte little-endian integers x_i and subtract a key k (deterministic, 16-bit or 24-bit found by heuristic) modulo 2^{24} :

$$x'_i = (x_i - k) \bmod 2^{24}.$$

The inverse adds back k . This reduces magnitude, making values more compressible by subsequent entropy coders.

- DOCX Paragraph and Table Transforms (Transforms 31 and 32). These use the python-docx library to extract text runs from paragraphs or tables, apply dictionary-based encoding to each run's text, and store the formatting attributes (font size, bold, italic, etc.) in a compact binary format. The original DOCX structure is reconstructed during decompression, ensuring perfect losslessness.

9. Self-Test and Exhaustive Verification

To confirm implementational correctness, PJP-6 performs an exhaustive self-test:

- For every single transform and every byte value, the forward transform is applied to the single-byte input, then the inverse is applied, and the result is compared with the original.
- For every pair sequence (2,704 pairs) and every byte value, the same round-trip test is conducted.
- A random 1,000-byte block is processed through a full compress-decompress cycle in both marker-free and safe modes.
- The empty input is tested in both modes.
- Dictionary tokenisation and detokenisation are verified on sample texts.
- Transforms 27–32 are tested on representative inputs.

All tests pass on all paths, demonstrating the mathematical invertibility of every component and the robustness of the auto-correction mechanism. This confirms that the system will never emit a corrupted file.

10. Conclusion

PJP-6 demonstrates that a data-driven approach – building a large space of mathematically invertible byte-level transforms, dictionary-based tokenisations, and quantum-inspired permutations, and exhaustively searching over them – can significantly improve compression ratios over standard algorithms. Every component is defined by an explicit formula, every transform is proven bijective, and the marker-free design with built-in self-verification and automatic safe fallback guarantees zero-error decompression. The system integrates feature engineering, model selection, and rigorous validation into a single pipeline. With 2,961 core transform paths, hybrid dictionary modes, and specialised transforms for DOCX, PJP-6 offers a flexible, general-purpose preprocessing layer that adapts to any input. The dual-mode backend with transparent self-correction makes the compressor fully self-sufficient and 100% lossless under all inputs. PJP-6 stands as a substantial contribution to lossless compression and a notable data-science project.